

PKL-Bench: A Multi-Tier Benchmark Dataset and Evaluation Suite for Asymmetric Game Dynamics in the Pro Kabaddi League

Ritesh Dobhal*

University of Delhi, Delhi, India

ORCID: [0009-0009-2773-4972](https://orcid.org/0009-0009-2773-4972)

October 2026

Abstract

While sports analytics benchmarks such as SoccerNet and NBA Play-by-Play have accelerated machine learning in symmetric invasion and continuous-transition games, contact invasion sports governed by asymmetric pursuit-evasion dynamics remain underrepresented in computational literature. In this paper, we introduce **PKL-Bench**, the first comprehensive, research-grade benchmark dataset and evaluation harness for professional Kabaddi. Spanning all 10 complete seasons (2014–2024) of the Pro Kabaddi League (PKL), PKL-Bench harmonizes 1,060 matches, 26,760 player boxscores, and 103,176 timestamped, discrete raid events into a standardized 5-tier relational schema distributed in both Apache Parquet and Frictionless CSV formats. We formulate Kabaddi as an asymmetric, discrete-time Markov Decision Process (MDP) exhibiting discrete structural phase transitions—including the bonus-line deactivation threshold ($N_{def} = 6$) and the super tackle regime ($N_{def} \leq 3$). We establish a strict mathematical score conservation theorem, proving 100.0% accounting adherence across all 1,060 matches, and implement an out-of-sample temporal evaluation protocol (Train: Seasons 1–8, Validation: Season 9, Test: Season 10) that prevents lookahead leakage. Finally, we provide standardized baseline implementations and leaderboards across four foundational tasks: discrete raid outcome prediction, dynamic in-game win probability modeling, pre-match spread forecasting, and context-adjusted player valuation. PKL-Bench is openly released under CC-BY-4.0 accompanied by a complete Gebru-compliant Datasheet for Datasets.

1 Introduction

Sports analytics has emerged as a fertile testbed for artificial intelligence, driving advances in multi-agent reinforcement learning (MARL), spatio-temporal tracking, and dynamic win probability modeling. However, modern research is overwhelmingly concentrated in association football [1, 2], basketball [3], and cricket. These sports share structural symmetries: teams field equal numbers of players who freely transition between attack and defense on open playing surfaces.

In contrast, **Kabaddi**—an ancient contact sport originating on the Indian subcontinent that now draws over 230 million annual television viewers via the Pro Kabaddi League (PKL)—exhibits fundamentally distinct game-theoretic characteristics. Kabaddi is an **asymmetric, sequential zero-sum pursuit-evasion game** [4]. In each raid, exactly one attacker (the raider) enters the opponent’s territory to confront an organized defensive unit of $N_{def} \in \{1, \dots, 7\}$

*Correspondence to: ritesh2005dobhal@gmail.com (University of Delhi). Benchmark repository: <https://github.com/Prompt-0/pkl-bench>

players within a strict 30-second shot clock. The sport features rule-induced non-linearities: when the defense drops to ≤ 3 players, the payoff for a successful tackle doubles (Super Tackle); when two consecutive raids yield no points, the third raid triggers mandatory offensive execution (Do-or-Die).

Despite its strategic depth, Kabaddi analytics has remained fragmented. Prior studies were restricted to isolated hackathons, proprietary sports API paywalls, or post-hoc aggregated network analyses [5, 6]. No unified, multi-season, play-by-play benchmark dataset with rigorous evaluation protocols has existed.

To resolve this deficiency, we present **PKL-Bench**, an open-science benchmark suite designed according to FAIR principles:

1. **Complete Historical Corpus:** Covering 10 seasons (2014–2024), 1,060 matches, 26,760 player boxscores, and 103,176 discrete play-by-play raid events.
2. **Mathematical Rigor:** Formalization of Kabaddi game dynamics as an MDP, backed by a programmatic score conservation audit proving 100.0% point accounting integrity.
3. **Strict Leakage-Free Temporal Splits:** Partitioning matches chronologically (Train: S1–S8, Validation: S9, Test: S10) to prevent lookahead bias.
4. **Four Standard Benchmark Tasks:** Clear loss functions, evaluation metrics, and verified baselines for raid outcome prediction, in-game win probability, match outcome forecasting, and player impact ratings.
5. **Turnkey Open Source Software:** An installable Python SDK (`pk1-bench`), Frictionless Data package, and complete Datasheet for Datasets [7].

2 Formal Formulation of Kabaddi Game Dynamics

2.1 Court Geometry and Spatial Coordinates

A regulation PKL mat measures 13.0 m $\times$ 10.0 m, bisected into two 6.5 m $\times$ 10.0 m half-courts by the Midline. Each half contains:

- **Baulk Line (3.75 m from Midline):** The minimum threshold a raider must cross for a raid to be legally valid.
- **Bonus Line (4.75 m from Midline):** Active strictly when $N_{def} \geq 6$. Planting a foot past the bonus line with the trailing foot airborne yields 1 point.
- **Lobbies (1.0 m border strips):** Out of bounds until physical touch is initiated between raider and defender, after which they expand the legal playing area.

2.2 The Kabaddi MDP

We formalize a Kabaddi match as a discrete-time Markov Decision Process over sequential raids $t \in \{1, \dots, T\}$. At raid t , the state is defined as:

$$S_t = \langle \Delta_t, \tau_t, h_t, N_{def,t}, D_t, B_t, r_t, \mathcal{D}_t \rangle \quad (1)$$

where $\Delta_t = \text{Score}_{attack,t} - \text{Score}_{defense,t}$ is the score differential, $\tau_t \in [0, 1200]$ is the clock seconds remaining in the current half, $h_t \in \{1, 2\}$ is the match half, $N_{def,t} \in \{1, \dots, 7\}$ is the number of active defenders, $D_t \in \{0, 1\}$ is the binary Do-or-Die indicator, $B_t = \mathbb{I}(N_{def,t} \geq 6)$ is the bonus line eligibility flag, $r_t \in \mathcal{P}$ is the raider identity, and $\mathcal{D}_t \subset \mathcal{P}$ ($|\mathcal{D}_t| = N_{def,t}$) is the active defensive squad.




Figure 1: **Strategic Phase Transitions in Kabaddi.** (Left) Empirical distribution of raid outcome categories across 103,176 events in PKL-Bench. (Right) Raider strike rate and super tackle probability as a function of defender count N_{def} , highlighting the Bonus Line activation threshold ($N_{def} \geq 6$) and the Super Tackle regime ($N_{def} \leq 3$).

$$\mathcal{Y} = \{\text{EMPTY, SUCCESSFUL_TOUCH, BONUS_ONLY, SUPER_RAID, UNSUCCESSFUL_TACKLE, SUPER_TACKLE}\} \quad (2)$$

2.3 Phase Transitions and Discontinuous Payoffs

As illustrated in Figure 1, Kabaddi exhibits two distinct structural phase transitions:

1. **The Super Tackle Regime** ($N_{def} \leq 3$): Defensive reward jumps from 1 point to 2 points (+100%), while raider risk remains fixed at 1 dismissal. This alters the defensive risk appetite, inducing cooperative chain traps.
2. **The Bonus Threshold** ($N_{def} = 6$): As the defense transitions from 6 to 5 players, the bonus line deactivates ($B_t = 0$). The raider can no longer earn an uncontested point without initiating physical contact, forcing aggressive penetrative maneuvers.

3 Dataset Architecture & Integrity Audits

3.1 Multi-Tier Relational Hierarchy

PKL-Bench structures data into five relational tiers:

- **Tier 0 (Metadata):** Seasons ($N = 10$), Teams ($N = 12$), Venues ($N = 28$), and Rulesets ($N = 1$).
- **Tier 1 (Players):** Master directory of 805 athletes with standardized roles, positions, and career totals.
- **Tier 2 (Matches):** 1,060 match summary records with final scores, half scores, toss outcomes, and margin metrics.
- **Tier 3 (Player Matches):** 26,760 player boxscores detailing raids, tackles, points, cards, super 10s, and high 5s.

Table 1: PKL-Bench Dataset Tier Statistics Across Seasons 1–10.

Tier	Entity	Instance Count	Parquet Size	CSV Size
Tier 0	Seasons	10	0.01 MB	0.00 MB
Tier 0	Teams	12	0.00 MB	0.00 MB
Tier 0	Venues	28	0.00 MB	0.00 MB
Tier 0	Rulesets	1	0.00 MB	0.00 MB
Tier 1	Master Players	805	0.03 MB	0.04 MB
Tier 2	Matches	1,060	0.06 MB	0.22 MB
Tier 3	Player Match Boxscores	26,760	0.29 MB	2.55 MB
Tier 4	Play-by-Play Raid Events	103,176	1.31 MB	12.29 MB
Total	Complete Corpus	131,852	1.70 MB	15.10 MB




Figure 2: **Empirical Scoring Dynamics Across 1,060 Matches.** (Left) Total combined match points distribution ($\mu = 67.2$ pts). (Right) Margin of victory distribution, demonstrating that over 32% of PKL fixtures are decided by ≤ 5 points.

- **Tier 4 (Play-by-Play Raids):** 103,176 granular raid records tracking chronological sequences, clock times, score margins, and categorical outcomes.

3.2 Score Conservation Theorem and Verification

Theorem 1 (Conservation of Kabaddi Scoring). *In any regulation Kabaddi match, for each team $i \in \{1, 2\}$, the total score S_i is an exact linear sum of disjoint point sources:*

$$S_i = \text{RaidPoints}_i + \text{TacklePoints}_i + \text{AllOutPoints}_i + \text{ExtraPoints}_i \quad (3)$$

where $\text{RaidPoints}_i = \text{Touch}_i + \text{Bonus}_i$, $\text{TacklePoints}_i = \text{Capture}_i + \text{SuperTackleBonus}_i$, and $\text{AllOutPoints}_i = 2 \times \text{AllOutsInflicted}_i$.

Empirical Verification: We executed automated verification across all 1,060 matches using `pkl_bench.validator`. The conservation theorem holds with **100.0% adherence** on both Team 1 and Team 2, with 0 discrepancies flagged across the 10-season corpus.




Figure 3: **In-Game Win Probability Dynamics (Task 2).** (Left) Reliability calibration curves on the out-of-sample Season 10 test set. (Right) Sample win probability trajectories comparing a tight comeback thriller against a decisive blowout.

4 Benchmark Tasks & Evaluation Framework

4.1 Strict Temporal Split Protocol

Random cross-validation creates severe data leakage in sports because player rosters, coaching staffs, and team momentum persist across fixtures. PKL-Bench enforces a strict chronological partition:

- **Train Set (Seasons 1–8, 2014–2022):** 787 matches (74.2%), 75,777 raid events.
- **Validation Set (Season 9, 2022):** 137 matches (12.9%), 13,505 raid events.
- **Test Set (Season 10, 2023–2024):** 136 matches (12.8%), 13,894 raid events.

Splits are mutually disjoint ($Train \cap Val = \emptyset$, $Train \cap Test = \emptyset$, $Val \cap Test = \emptyset$).

4.2 Task Definitions

- **Task 1: Discrete Raid Outcome Prediction:** Multi-class classification predicting $y_t \in \mathcal{Y}$ from pre-raid state features. Evaluated by Multiclass Log-Loss and Macro-F1.
- **Task 2: Dynamic Win Probability:** Estimating win probability $P(W_1 | S_t)$ at every raid t . Evaluated by Brier Score [8] and Expected Calibration Error (ECE) [9].
- **Task 3: Match Forecasting:** Predicting match winner and point spread prior to coin toss. Evaluated by Accuracy, ROC-AUC, and Margin MAE.
- **Task 4: Player Valuation:** Context-conditioned Expected Points Added (EPA) adapting action-value valuation paradigms [10] and True Raider Impact (TRI).

5 Experimental Results & Leaderboard

Table 2 presents baseline benchmark results evaluated on the unseen Season 10 test set.

Key Insights:




Figure 4: **Franchise Elo Trajectories Across 10 Seasons (Task 3).** Dynamic ratings of leading PKL franchises from Season 1 to Season 10, capturing historical dynasty runs (e.g., Patna Pirates’ three-peat in Seasons 3–5).

Table 2: **Official PKL-Bench Baseline Leaderboard (Evaluated on Season 10 Test Set).**

Task	Model	Primary Metric	Secondary Metric	Tertiary Metric
Task 1: Raid Outcome	Majority Class	Accuracy: 0.4770	Macro-F1: 0.1292	Log-Loss: 18.8520
	Multinomial Logistic	Accuracy: 0.5188	Macro-F1: 0.2274	Log-Loss: 1.0626
	HistGradientBoosting	Accuracy: 0.5442	Macro-F1: 0.2484	Log-Loss: 0.9597
Task 2: Win Probability	Logistic Leverage	Brier: 0.2270	ECE: 0.1882	Log-Loss: 0.6446
	Calibrated GBDT	Brier: 0.1991	ECE: 0.1561	Log-Loss: 0.5827
Task 3: Match Forecasting	Random Baseline	Accuracy: 0.5000	Brier: 0.2500	ROC-AUC: 0.5000
	Dynamic Kabaddi Elo	Accuracy: 0.7059	Brier: 0.1968	ROC-AUC: 0.7707
Task 4: Player Valuation	Raw Total Points	Top Volume: P. Narwal (1,690)	Top Tackle: F. Atrachali (486)	Naive
	Expected Points Added	Top Cum.: P. Narwal (+470.3)	Top Rate: P. Sehrawat (0.190)	State-Conditioned

- On Task 1, Gradient Boosting achieves a substantial log-loss reduction from 1.0626 to 0.9597 over Logistic Regression, driven by non-linear decision trees capturing the interaction between Do-or-Die flags and clock urgency.
- On Task 2, Calibrated GBDT reduces Expected Calibration Error (ECE) from 0.1882 to 0.1561, producing reliable probability outputs (Figure 3).
- On Task 3, Dynamic Kabaddi Elo with margin-of-victory scaling achieves 70.6% accuracy and 0.7707 ROC-AUC on out-of-sample matches, confirming that franchise strength exhibits meaningful temporal autocorrelation across 10 seasons (Figure 4).

6 Conclusion & Open Research Directions

PKL-Bench establishes the first standardized, mathematically verified benchmark dataset for professional Kabaddi, spanning 10 complete seasons and 103,176 raid events. By releasing this resource under CC-BY-4.0 with reproducible baselines and strict temporal splits, we provide the computational research community with a fertile testbed for asymmetric game theory, multi-agent reinforcement learning, and calibrated sports analytics.

Future work will expand PKL-Bench with multi-camera computer vision pose tracking and continuous spatial player trajectory tracking.

References

- [1] A. Delière, A. Cioppa, S. Giancola, M. J. Seikavandi, J. V. Dueholm, K. Nasrollahi, B. Ghanem, T. B. Moeslund, and M. Van Droogenbroeck, “Soccernet-v2: A dataset and benchmarks for holistic understanding of broadcast soccer videos,” *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4508–4519, 2021.
- [2] L. Pappalardo, P. Cintia, A. Rossi, E. Massucco, P. Ferragina, D. Pedreschi, and F. Giannotti, “A public data set of spatio-temporal match events in soccer competitions,” *Scientific Data*, vol. 6, no. 1, pp. 1–15, 2019.
- [3] A. Franks, A. Miller, L. Bornn, and K. Goldsberry, “Characterizing the spatial structure of defensive skill in professional basketball,” *The Annals of Applied Statistics*, vol. 9, no. 1, pp. 94–121, 2015.
- [4] R. Isaacs, *Differential games: a mathematical theory with applications to warfare and pursuit, control and optimization*. Courier Corporation, 1999.
- [5] A. Sengupta, S. Layek, and K. Deyasi, “Analyzing player involvement in the indian pro kabaddi league: A network analysis approach,” *arXiv preprint arXiv:2405.01596*, 2024.
- [6] B. Lalwani and A. Mukherjee, “Kabaddipy: A package to enable access to professional kabaddi data,” *arXiv preprint arXiv:2501.05168*, 2024.
- [7] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford, “Datasheets for datasets,” *Communications of the ACM*, vol. 64, no. 12, pp. 86–92, 2021.
- [8] G. W. Brier, “Verification of forecasts expressed in terms of probability,” *Monthly Weather Review*, vol. 78, no. 1, pp. 1–3, 1950.
- [9] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *International Conference on Machine Learning*. PMLR, 2017, pp. 1321–1330.
- [10] T. Decroos, L. Bransen, J. Van Haaren, and J. Davis, “Actions speak louder than goals: Valuing player actions in soccer,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2019, pp. 1851–1861.